



1. Parallel Programming - the Context

Parallel Programming
Dr Hamidreza Khaleghzadeh
School of Computing
University of Portsmouth



Goals

- Set the context for learning parallel programming
 - the emergence of *multicore* and *many-core* commodity systems
 - the nature of modern *supercomputers*.
- Structure of module and assessment.



Parallel Computing

- This module will focus on parallel computing and especially parallel programming.
- We interpret “parallel computing” as:
Harnessing the nodes of multi-core or many-core processors, or computer clusters including supercomputers, to work cooperatively on individual problems, with a goal of getting results fast.



Context

- Until quite recently parallel computing has been the province of researchers with access to *supercomputers* - e.g. in university departments, national labs, weather forecasting centres, etc.
- The situation has changed in recent years with the widespread availability of *multi-core processors* and *many-core* accelerators.
- Let's reprise some of that context
 - If you attended the final TB1 lecture of COSINE in the second year, some of this may look familiar...



MULTICORE AND MANYCORE



The Era of the Free Lunch


- Increasing *clock speeds* (combined with increasing *Instruction Level Parallelism*) over the decades leading up to the early 2000s meant that programs would run faster and faster with no extra effort
- ... or, equivalently, that applications of greater and greater complexity could be executed on hardware of fixed or reducing cost.



Concerning Processor Speed

- *Clock speed* of microprocessors increased exponentially through the 1990s and beyond; but for now they appear to have reached a limit.
- In 2005 Herb Sutter pointed out that clock speed for Intel CPUs had *stuck* around 3.5 GHz.
 - ◉ Fifteen years on, observation still good.
 - ◉ Clock speed is now limited by power consumption/heat dissipation.
- Moore's Law for *transistor count* continues (for now).


Intel CPU Clock Speed (to 2009)[†]



[†] Graph taken from: <http://www.gotw.ca/publications/concurrency-ddj.htm>

Our World
in Data

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important as other aspects of technological progress – such as processing speed or the price of electronic products – are linked to Moore's law.



Data source: Wikipedia (https://en.wikipedia.org/wiki/Transistor_count)
The data visualization is available at [OurWorldinData.org](https://ourworldindata.org). There you find more visualizations and research on this topic.


Licensed under [CC-BY-SA](#) by the author Max Roser.



Multicore


- Response of chip manufacturers has been *multicore* - multiple independent CPUs on a single chip.
 - Already commonplace - all but cheapest new desktops quad-core; laptops at least dual-core.
 - 8-core Nehalem-EX chip (for servers) released in March 2010:
http://www.theregister.co.uk/2010/03/30/intel_nehalem_ex_launch/
 - 15-core Intel Xeons released early 2014:
<http://techreport.com/news/26056/intel-releases-15-core-xeon-e7-v2-processor>
 - 28-core Xeons July 2017
 - (56-core available April 2019, but dual processor)

Example Quad-core Processor†



†Intel Core i7 "Nehalem", from: <http://www.hardwarecanucks.com>

Intel Xeon Maximum Cores by Year[†]



[†]Data from:

https://en.wikipedia.org/wiki/List_of_Intel_Xeon_microprocessors



Graphics Processing Units

- GPUs are widely available *Massively Parallel Processors*.
- Rather than a few complex, superscalar cores, they use *many, simple* cores.
 - E.g. NVIDIA GTX 680:
<http://developer.nvidia.com/content/geforce-gtx-680-developers-perspective>
has 1536 "CUDA" cores.
- Notional performance for floating point: *3090 GFLOPS*
 - c.f. perhaps few 10s GFLOPS for a multicore *CPU*.



"General Purpose" GPUs

- Using APIs like *CUDA* and *OpenCL*, it is possible to use GPUs for general processing (not just graphics).
- Benchmark on next slide from machine identical to the PCs in BK0.19.



Running NVIDIA Sample code for Matrix Multiplication

[Matrix Multiply CUBLAS] - Starting...

GPU Device 0: "GeForce RTX 2070" with compute capability 7.5

GPU Device 0: "GeForce RTX 2070" with compute capability 7.5

MatrixA(640,480), MatrixB(480,320), MatrixC(640,320)

Computing result using CUBLAS...done.

Performance= 3650.41 GFlop/s, Time= 0.054 msec, Size= 196608000 Ops

Computing result using host CPU...done.

Comparing CUBLAS Matrix Multiply with CPU results: PASS

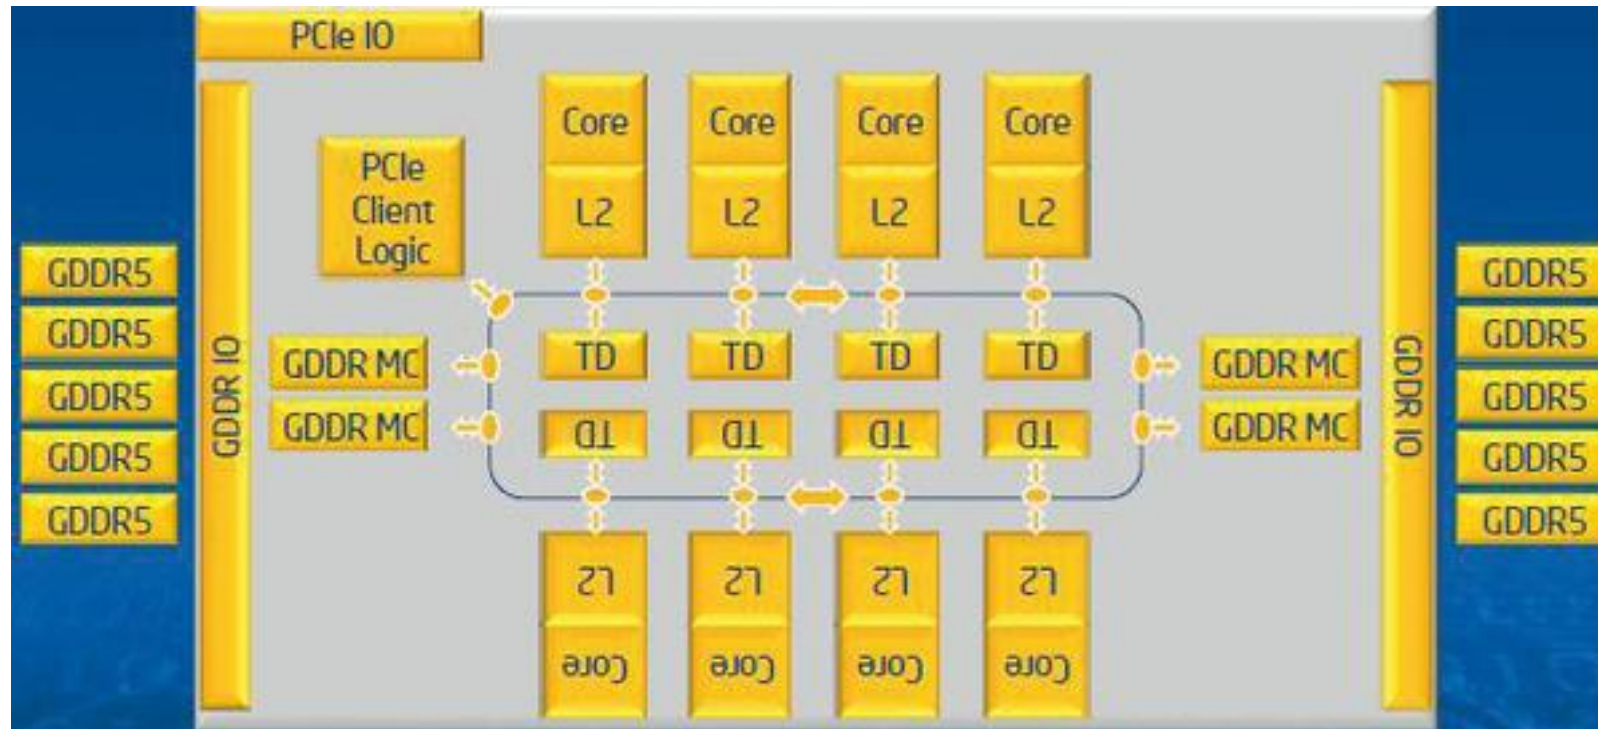
NOTE: The CUDA Samples are not meant for performance measurements.
Results may vary when GPU Boost is enabled.



Xeon Phi

- An alternative to the GPU co-processor approach, being marketed by Intel, is their *Multiple Integrated Core Architecture (MIC)*.
- In this product line, a large number of Pentium-based cores are integrated on a single chip.
 - Unlike GPU, each core is a *general x86 processor* with private level 1 and 2 cache.
 - Unlike conventional multicore CPU, there is no *shared* cache on chip, but there is cache coherent access to external memory.
 - Fairly conventional programming tools apply.

MIC Architecture[†]



- Current generation Xeon Phis have *up to 72 cores* on chip.

[†]From: <http://www.admin-magazine.com/Articles/Exploring-the-Xeon-Phi>



Software Implications

- *But, as Sutter said, move to multicore and many cores marks “the end of the free lunch”.*
 - Applications no longer *automatically* speed up with increasing clock rate.
- To exploit multiple cores, apps must be *rewritten*, adopting *parallel programming* techniques.
- This may eventually impact much or all of the software industry!



SUPERCOMPUTERS



Supercomputers and Parallelism

- Supercomputers have been exploiting massive parallelism for some decades, and many parallel programming techniques were developed for these.
- Increasing numbers of cores in commodity chips does *not* mean the need for supercomputers has gone away!
- In fact most supercomputers today are built by bolting together *large* numbers of “ordinary” multicore CPUs, to get even more parallelism.



Benchmarking Supercomputers

- Many different benchmarks suites have been developed to measure the speed of supercomputers.
- One of the simplest and most widely used numerical benchmarks today is called *HPL*[†].
- In HPL, one solves a system of linear equations:

$$A x = b$$

A is an N by N random matrix, b is a random vector, and x a vector that must be solved for.

[†]<http://www.netlib.org/benchmark/hpl/>

High-Performance Linpack Benchmark(HPL)

- Solve for all x_i :

$$A_{11}x_1 + A_{12}x_2 + \dots + A_{1N}x_N = b_1$$

$$A_{21}x_1 + A_{22}x_2 + \dots + A_{2N}x_N = b_2$$

... ..

$$A_{N1}x_1 + A_{N2}x_2 + \dots + A_{NN}x_N = b_N$$

- A and b are randomly generated.
- The solution must be done by an LU scheme involving $2/3 N^3 + O(N^2)$ floating point operations.
- The benchmarker must adjust N to get the *best performance* in FLOPS for the system being tested.



FLOPS Reminder

- Number of *floating point operations* executed *per second* (1 FLOP = a single +, -, *, etc operation)
 - *megaFLOPS* ($= 10^6$ FLOPS) - historic processors
 - *gigaFLOPS* ($= 10^9$ FLOPS) - current CPUs
 - *teraFLOPS* ($= 10^{12}$ FLOPS) - current GPUs
 - *petaFLOPS* ($= 10^{15}$ FLOPS) - world class, massively parallel supercomputers
 - *exaFLOPS* ($= 10^{18}$ FLOPS) - not quite yet - see later!



TOP500

- The *TOP500 list* is published twice a year[†]
- Often regarded as a ranking of the most powerful supercomputers in the world
- Technically speaking, it only ranks their speed at executing HPL!
 - There seems to be some move afoot to shift emphasis to a new benchmark High Performance Conjugate Gradients, claimed to be representative of a wider range of real applications.

[†]<http://www.top500.org/>



Sierra

- *Number 5* on the June 2022 TOP500 list.
- Located at Lawrence Livermore National Lab in US.
 - Commissioned for nuclear weapons simulation.
 - LINPACK performance of *94.64 petaFLOPS*.
 - Consists 8,640 of 22-core *POWER9* processors for (only!) *190,080 cores*
 - ...connected to 17,280 NVIDIA V100 GPUs.



Summit

- *Number 4* on the June 2022 TOP500 list.
- A *IBM* developed computer at Oak Ridge National Lab in US.
 - For general use by scientific community - initial apps likely to include cosmology, medicine and climatology.
 - LINPACK performance of *143.5 petaFLOPS*.
 - Consists 9,216 of 22-core POWER9 processors for *202,752 CPU cores*
 - *But* POWER9s connected to 27,648 Nvidia Tesla V100 GPUs, each with 5120 CUDA cores - so notional total of *141,557,760 CUDA cores!*



Supercomputer Fugaku

- *Number 2* on the June 2022 TOP500 list.
- Developed at RIKEN Centre for Computational Science, Kobe Japan.
 - Ultimate goal to break the *exaFLOP* barrier.
 - LINPACK performance of *415.5 petaFLOPS*.
 - Consists 158,976 Fujitsu A64FX processors (48-core *ARM*-based processor, with *Scalar Vector Extensions* SIMD capabilities) for *7,640,848 CPU cores*.



Supercomputer Frontier

- *Number 1* on the June 2022 TOP500 list.
- Developed at Hewlett Packard Enterprise (HPE)
 - LINPACK performance of *1,102 petaFLOPS*.
 - Consists 136,408 AMD Optimized 3rd Generation EPYC 64C 2GHz for *8,730,112 CPU cores*.



Supercomputing at UP

- The university's Institute of Cosmology and Gravitation run a smaller supercomputer called *SCIAMA* (*SEPNet Computer Infrastructure for Astrophysical Modelling and Analysis*)[†]
 - Primarily for cosmological simulation and data analysis, but available for researchers throughout the university.
 - Originally assembled from 6-core Intel Xeon Westmere processors for *1008 CPU cores*.
 - Recently upgraded to *SCIAMA II*, adding Ivy Bridge processors to more than double total number of cores.

[†] <http://www.port.ac.uk/research/news/title,149485,en.htm>
<http://www.sciama.icg.port.ac.uk/>



The Millennium Simulation

- Slightly old now, but nice example of a supercomputer application in cosmology
 - See *Simulating the Joint Evolution of Quasars, Galaxies and their Large-Scale Distribution* at <http://www.mpa-garching.mpg.de/gadget>
 - Also see [movies](#) at that site!
- Follows evolution of 10^{10} dark matter “particles” from early Universe ($z = 127$) to current day.
 - Performed on 512 nodes of IBM p690.
 - Used 1TB of distributed memory.
 - 350,000 CPU hours - 28 days elapsed time.
 - Floating point performance around 0.2 TFLOPS.



PLAN FOR THIS TEACHING BLOCK



Module Structure

- General plan:
- Learning driven by *practical work*; assessment takes form of a “lab book”.
- Early lab sessions will be “scripted”, and introduce some basic skills for parallel programming.
- Some later lab sessions will be dedicated to an individual development project that you will select about half-way through the term.



Scripted Lab Sessions

- These will use Java, starting with some simple examples of parallel programming using Java threads on the multi-core machines in the labs.
- Subsequent scripted sessions will use an *MPI*-like Java interface called *MPJ Express* to introduce parallel programming on “clusters” (hopefully using subsets of machines in the lab as our clusters).



Development Project

- I will provide a few possible choices of programming project, or you can propose your own (I must agree to it!)
 - You may use Java/MPJ Express for implementation in the teaching labs.
 - If you wish to use C/C++ or other programming environments. *Should* be possible in labs using CentOS - up to you to assess feasibility!
 - Or you might bring your own platform - perhaps you have an NVIDIA GPU in your laptop that you can use for CUDA programming?!



Lab Books

- 4000 words, which should record all lab work, and the outcome of your development project.
- Details of marking scheme, etc, will follow shortly. For now, make sure you keep a record of all you do in lab sessions!



Lectures

- Early lectures concentrate on introducing programming environments, paradigms and algorithms that are likely to be useful in your lab work (including development project), with an emphasis on MPI programming.
- Later lectures will cover more diverse topics, including more novel approaches to programming parallel systems, with recent developments.



Summary

- Set the context for parallel programming and this unit.
- Next week: *parallel programming with threads*



Further Reading

- See links embedded in these slides.