



2. Decomposing the Problem

Parallel Programming

Dr Hamidreza Khaleghzadeh
School of Computing
University of Portsmouth



Goals


- Start with a general overview of programming frameworks for parallel computing, contrasting *shared memory* and *distributed memory*.
- Then, with reference to a simple example, discuss issues of problem decomposition and load balancing.



PARALLEL PROGRAMMING

Parallel Computer Organization

- Traditional dichotomy in parallel computing - shared vs distributed memory (*SMP*[†] vs *cluster*):



[†]*Symmetric Multi-Processor.*



SMP Programming Frameworks

- Shared memory (cooperating *threads*):
 - POSIX or Java Threads
 - OpenMP
 - Cilk
 - Intel Threaded Building Blocks
 - etc, etc



Threads

- In C/C++ can use the *POSIX threads*[†] (*pthread*s) libraries to achieve concurrency.
- In Java, use native *Thread* class for same effect - we are using this approach in early labs.
- On *multicore* system, threads automatically scheduled on different cores to give parallel execution.

[†] There is a decent tutorial at <http://www.yolinux.com/TUTORIALS/LinuxTutorialPosixThreads.html>, but note pthreads is not limited to Linux – also works on Windows, for example!



OpenMP

- *OpenMP*[†] (which rather misleadingly styles itself an “API”) is a set of *compiler directives* and supporting libraries for exploiting *shared memory* computers.
- Defined for C/C++ (or Fortran)
- Parallelizing a loop in OpenMP *may* be as simple as this:

```
#pragma omp parallel for  
for(i = 0 ; i < N ; i++) {  
    ...  
}
```

[†]www.openmp.org



Cluster Programming Frameworks

- Distributed memory (cooperating *processes*)
 - MPI (PVM, etc)
 - Co-array Fortran, UPC, etc
 - Global Array Toolkit (etc)
 - etc, etc




MPI

- *MPI*[†] (simply *Message Passing Interface*) is a standard programming interface for sending and receiving messages in C or Fortran programs.
- Beyond this, provides most popular framework for *Single Program Multiple Data* (SPMD) computing on large scale distributed memory parallel computers - i.e. *clusters*.
- In labs we will use a Java version of MPI called *MPJ Express*.

[†]<http://www-unix.mcs.anl.gov/mpi>

Real World - Hybrids

- Clusters of shared memory nodes - typical of large modern systems
- e.g. Portsmouth SCIAMA supercomputer





Programming Hybrid Machines

- Can be programmed using a combination of shared memory framework (e.g. threads or OpenMP) *within* nodes, plus distributed memory framework (e.g. MPI) *between* nodes.
- Sometimes easier to adopt the lowest common denominator - use MPI (say) across *all* cores. But efficiency lower.



SMPs - Pros and Cons

- Parallel programming is generally easier for shared memory systems - also called *Symmetric Multi-Processors* (SMPs).
 - No need to split up data structures into separate pieces held in memory of different nodes
 - No need for explicit communication operations between nodes - inter-node communication happens implicitly through memory access (usually *synchronization* operations will still be required).
- Commodity multicore microprocessors are generally SMPs.
 - But the *largest* parallel systems are *not*.

A Large SMP[†]



- COSMOS supercomputer in Cambridge for astrophysics, launched (by Stephen Hawking) 2012
- One of the biggest *shared memory* computers in Europe. SGI Altix UV2000 built from Intel Xeon processors with total **1856** cores.
- But recall *largest* supercomputers have *millions* of cores!

[†]https://safe.epcc.ed.ac.uk/diracwiki/index.php/Cambridge_COSMOS_SHARED_MEMORY_Service



Early Lectures

- Shared memory systems are ultimately less scalable than distributed memory.
- But, because easier, we begin our exploration of parallel programming assuming shared memory systems.
- Later lectures will explore parallel programming for clusters.




PROBLEM DECOMPOSITION



First Example


- In the labs you should already have started looking at parallelizing the *Mandelbrot Set*.
- Good starting example because fun and easy to parallelize ("embarrassingly parallel").
- But also illustrates significant issues about *problem decomposition* - always important in parallel programming.

Mandelbrot Set



Mathematically,
the set is the
black area within
part of the x, y
(or complex
number) plane
with $-2 \leq x, y \leq 2$

Mandelbrot Set (Maths)



Define complex:

$$c = x + iy$$


and recurrence:

$$z_0 = c$$

$$z_{k+1} = c + z_k^2$$

If $\text{abs}(z_k) < 2$ for
all $k < \infty$, then
point x, y is in the
set (black area).

Mandelbrot Set (Numeric)



Consider only x, y values in N by N grid ("pixels").

If $\text{abs}(z_k) < 2$ for all $k < \text{CUTOFF}$, then point x, y is in the set (black area), where CUTOFF is some sufficiently large number.



Pseudocode

```
for(int i = 0 ; i < N ; i++) {  
    for(int j = 0 ; j < N ; j++) {  
  
        double x = step * i - 2.0 ;    //  $-2 \leq x \leq 2$   
        double y = step * j - 2.0 ;    //  $-2 \leq y \leq 2$   
  
        complex c = (x, y) , z = c ;  
  
        int k = 0 ;  
        while (k < CUTOFF && abs(z) < 2.0) {  
            z = c + z * z ;  
            k++ ;  
        }  
        set [i] [j] = k ;  
    }  
}
```



Loop Structure


- Don't worry about details of z recurrence here, but *do* look at general loop structure.
- Mathematically, *innermost while loop* repeats *forever* if we are in the black region; in practice, stop the loop after some number **CUTOFF** of iterations.
- *Upshot*: *inside or near* the black area there are many more iterations - and *more computational work* - than for typical points in the white area (where the while loop may end after a few iterations).



Parallel Decomposition

- If we want to calculate the set on, say, two cores, an obvious thing is to break the x,y plane into two halves - we may either do this vertically or horizontally.


Divide the i -loop



- Poor *load balancing*; core 0 has more work (more black!) than core 1.
- Work not evenly divided between cores, so imperfect parallel speed up.

Divide the j -loop

Core 0



Core 1



- Perfect *load balancing* because of symmetry of set.
- Parallel speed-up factor may be close to 2.



But...

Core 0

Core 1

Core 2


Core 3

- Disastrous *load balancing* - cores 0 and 3 only have about a tenth of the work of cores 1 and 2.
- Parallel speed-up on four cores not much better than on two cores!

Workload Distribution Formats

- Both horizontal and vertical decompositions are using a *block-wise* distributions of *index space*:

Sequential




Parallel

"Cyclic" Distribution


- An alternative distributions of *index space* from original problem:

Sequential



Parallel

Mandelbrot with Cyclic Decomposition



- Much more even break down of set.
- Good load balance.



Programming Considerations

- Assume original loop is
for (int i = 0 ; i < N ; i++) {...}
- Block-wise decomposition, each thread does[†]:
for (int i = me*N/P; i < (me+1)*N/P ; i++) {...}
- Cyclic decomposition, each thread does
for (int i = me ; i < N ; i+=P) {...}
- Here P is number of threads, and me is current thread identity.


[†]Formulas here get slightly more complicated if P doesn't divide N exactly.



Is Cyclic Decomposition Just Better?

- Not always, by any means.
- Even on shared memory computers, it sometimes leads to poor *cache* utilization, because of the stride in the local loop ($i += P$).
- On distributed memory computers, as we will see later, it often leads to *excessive communication requirements*.
- But for problems where *load balancing* is the main problem it can be very useful.

Other Approaches to Load-Balancing



- Divide the work into *many small chunks* - in this case many small tiles of Mandelbrot
- Need *many more tasks* than number of available *cores*.



Dynamic Load Balancing

- Each core grabs one tile, more or less randomly, and works on that.
- If it finishes before all tiles have been processed, it grabs another tile, and so.
- Keeps busy.
- Obviously more complex to program, but will discuss in more detail later.



ASIDE: PARALLEL SPEEDUP AND PARALLEL EFFICIENCY



Parallel Speedup

- One way of expressing the performance of a program is in *GigaFLOPS*. For parallel programs, another way is in terms of *speedup*.
- Suppose the execution time of the sequential version of the program is T_S , and the execution of a parallel version of the same program running on P cores is T_P .
- We define the *parallel speedup* as:
$$T_S / T_P$$
- What is the *ideal* parallel speedup on P cores?



Parallel Efficiency

- *Parallel efficiency* is the ratio of the *actual* speedup of a parallel program to the ideal speedup:

$$(T_S / T_P) / P$$

The efficiency is normally a number between 0 and 1.

- It measures how effectively we are making use of the P available cores.
- Parallel efficiency can fall due to poor load balancing, or other overheads of parallelism.

An Example from Above

Core 0

Core 1

Core 2

Core 3

- Assume the only inefficiency is due to load balancing...
- ... and cores 0 to 3 respectively do:
5%, 45%, 45%, 5%
of total work.



Speedup and Efficiency for Example

- Time for individual cores to complete work is:

$$0.05T_S, 0.45T_S, 0.45T_S, 0.05T_S$$

and parallel program as a whole finishes in *longest* of these times:

$$T_P = 0.45T_S$$

- *Parallel speedup* is:

$$T_S / T_P = T_S / (0.45T_S) = 1/0.45 \approx 2.22$$

- *Parallel efficiency* is

$$(T_S / T_P) / P = 2.22 / 4 \approx 0.55$$

or 55%.



ASIDE: BENCHMARKING PROTOCOL



Reproducibility of Timings

- Benchmarking parallel and sequential programs will be vital throughout the labs.
- For various reasons the time a program takes to execute is usually not completely reproducible - it varies from run to run.
- Where high quality timing results are needed it is advised to run the program several times (e.g. 4, 10, or more times) and record *all* the times.
 - To start with, this gives you a feeling for how variable the timing results are.
 - But then what should you record as the final result?



Summary Timings

- Many people's intuition is to take the *average* of all the runs as the best estimate of the time.
- I recommend that instead you take the *minimum* time of all the runs as the final result.
 - This may sound like cheating. But usually most variation from the minimum time is due to events external to the program itself (OS threads etc).
 - Using the minimum time usually leads to smoother and more reproducible graphs, etc.



Summary

- Next week: Topics in *shared memory parallel programming*.



Further Reading

- See links embedded in these slides.