



3. Shared Memory Parallel Programming

Parallel Programming

Dr Hamidreza Khaleghzadeh
School of Computing
University of Portsmouth



Goals

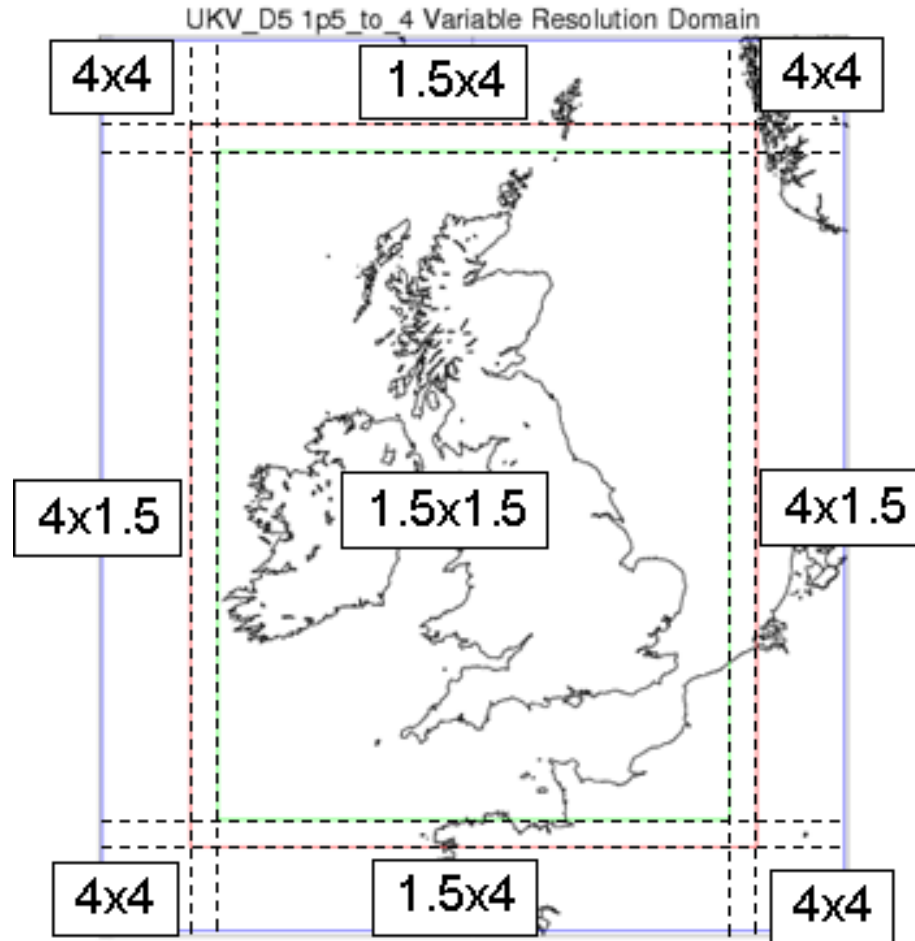
- Begin with general discussion of differences between “embarrassingly parallel” problems, and harder problems like simulations.
- In the context of the simple simulation introduced in this week’s lab script, introduce a fundamental kind of synchronization for shared-memory parallel programming.
- Brief overview of OpenMP - one of the established frameworks for shared memory programming.



Embarrassing Parallelism

- The Mandelbrot set was an example of an *Embarrassingly Parallel* problem, where each point is calculated independently of every other.
- Actually such problems are quite common - general situation is where a large problem decomposes into completely independent tasks.
- But many important problems *don't* have this property.


Weather Forecasting



- Size of individual cells (in kilometres) in Met Office's *UKV model*[†] of UK weather.

[†]<https://www.metoffice.gov.uk/research/approach/modelling-systems/unified-model/weather-forecasting>

Schematic of Grid





UKV Model

- Area of the UK is divided into 1.5km square cells horizontally, and into 70 layers vertically - a 3D grid (total $744 \times 928 \times 70$ grid points)
- Numerical modelling techniques used:
 - Atmospheric variables (wind speed, temperature, pressure, humidity, ...) stored for each grid box.
 - Equations are solved for each grid box to predict the values at that point a short time later.
 - Repeat until reach forecast time required.
 - For UKV, time step is 50 seconds, and forecast goes out to 36 hours.





Running Weather Simulation

- Dividing the grid over processors gives opportunity for exploiting massive parallelism
 - As of 2017, Met office runs a Cray XC40 supercomputer with total 241,920 cores that was #15 in Nov 2017 TOP500.
- But not “embarrassingly parallel” - equations at each grid point will depend on atmospheric variables stored at *neighbouring* points.
 - Each element of atmosphere moved by pressure of neighbouring elements, is warmed or cooled by them, etc.

Simplified Dynamics

- Net force on volume of air in *central* cell (tending to cause air to accelerate in x -direction), is proportional to difference in air pressures in *adjacent* cells, $P_1 - P_2$.





SIMULATIONS



Simpler Models

- In this unit we can't do weather forecasting.
 - Equations for updating variables within grid points involve complex physics that is out of scope.
- But we *can* look at simpler problems that illustrate the features affecting parallelization.
- The Conway's Game of Life example in this week's lab already incorporated some of these
 - Update of an individual cell according to some local rule, which however depends of values of neighbours.



"Life"

- Cellular automaton on 2D grid, where each cell takes value 0 or 1.
- In a time step, new value of cell depends on it's old value and old values of it's neighbouring cells.
- For current purposes, close enough analogy with large scale simulations on grids (even 3D) where local update depends on neighbours.

Life Animation[†]




[†]From https://www.conwaylife.com/wiki/Conway%27s_Game_of_Life

Update Rule for Cell

Sum of
neighbours = 3


1	0	0
0	1	1
0	0	1



	1	

Stay the same!


1	0	0
0	0	1
0	0	1



	0	

Sum of
neighbours = 2


0	1	0
0	1	0
1	0	0



	1	

Live!


0	1	0
0	0	0
1	0	0



	1	

Sum of
neighbours
anything else


1	0	0
0	1	1
1	1	0



	0	

Die!

1	0	0
0	0	1
1	1	0



	0	



Sequential Life

- Pseudocode - note two phases, "sum" and "update":

```
while(true) {  
    for(int i = 0 ; i < N ; i++)  
        for(int j = 0 ; j < N ; j++)  
            sums [i] [j] = sum of cells values neighbouring (i, j) ;  
    for(int i = 0 ; i < N ; i++)  
        for(int j = 0 ; j < N ; j++)  
            cells [i] [j] = update(cells [i] [j], sums [i] [j]) ;  
}
```

- Main array is **cells**; auxiliary array **sums** holds sum of cell elements neighbouring (i, j)




Aside

- Why do we need *two* sets of *for* loops here - why couldn't we have, say?:

```
while(true) {  
    for(int i = 0 ; i < N ; i++)  
        for(int j = 0 ; j < N ; j++) {  
            sum = sum of cells values neighbouring (i, j) ;  
            cells [i] [j] = update(cells [i] [j], sum) ;  
        }  
}
```

Updating cells array in iteration t

- Old cell value – after iteration $t - 1$
- New cell value – after iteration t



- Iterate through array, updating cells.
- With code on previous slide, **sum** for next (yellow) cell contains 4 green values and 4 red values.
- Wrong! Update of cell in iteration t should only depend on values of cells in **previous** generation.



Parallel Attempt

```
Class SumThread {  
    void run() {  
        for(int i = begin ; i < end ; i++)  
            for(int j = 0 ; j < N ; j++)  
                sums [i] [j] = sum of cells values neighbouring (i, j) ;  
    }  
}
```

```
Class UpdateThread {  
    void run() {  
        for(int i = begin ; i < end ; i++)  
            for(int j = 0 ; j < N ; j++)  
                cells [i] [j] = update(cells [i] [j], sums [i] [j]) ;  
    }  
}
```



Parallel Attempt (continued)

- Main program:

```
while(true) {  
    create P instances of SumThread, start them, and wait  
        for completion ;  
    create P instances of UpdateThread, start them, and  
        wait for completion ;  
}
```

- On previous slide we left *begin* and *end* undefined, but may assume a *block-wise decomposition* of *i* loops.
 - See previous lecture.



Comments

- The approach given over the previous two slides is clearly correct, and follows approaches taken in labs for parallelizing (i, j) loops in Mandelbrot Set.
- *But*, the *overhead of creating new threads* in every iteration of the main loop is likely to lead to *poor performance*
 - Inner (i, j) loops here typically smaller and contain much less work than Mandelbrot set.

Second Attempt

```
Class LifeThread {  
    void run() {  
  
        while(true) {  
            for(int i = begin ; i < end ; i++)  
                for(int j = 0 ; j < N ; j++)  
                    sums [i] [j] = sum of cells values neighbouring (i, j) ;  
            for(int i = begin ; i < end ; i++)  
                for(int j = 0 ; j < N ; j++)  
                    cells [i] [j] = update(cells [i] [j], sums [i] [j]) ;  
        }  
    }  
}
```



Second Attempt (continued)


- Main program:

create P instances of LifeThread, start them, and wait for completion ;

- This version creates the threads only once at the start of the program, and runs the whole of the main loop over time in each thread.
 - *Relatively* , much smaller overheads from thread creation - much more work in each thread.
- But is it correct??

A Problem

- The problem occurs for the cells at the edges of each thread's *domain* (cells it updates):





Shared Access

- The blue cell is updated by core 0. But the decision rule for its update depends on all the yellow cells.
- Hence core zero is *writing* the blue cell, and *reading* all the yellow cells.
 - In general the green cells, while written by the thread they “belong” to, are also read by the thread to their left.
 - Red cells are written by the “owning” thread, but also read by the thread to their right.
- These array elements are *shared variables*.



Outcome

- It is very likely that some threads will run ahead, maybe by several generations, of other threads.
- When a thread is updating cells on the edges of its domain in terms of cells updated by an adjacent thread, it may well be reading the wrong generation of cells.
- Effect is unpredictable and non-deterministic - it is an example of what kind of condition in a concurrent program?

Race condition.



A New Kind of Synchronization

- In the 2nd year OS unit when we encountered race conditions, a common solution was to employ the form of *thread synchronization* called *mutual exclusion*.
- Here we need a more general, global form of synchronization, that keeps all threads in *lock-step*.
- The mechanism we need is *barrier synchronization*.



BARRIERS



The Barrier

- We will regard a *barrier* as an object that P threads share access to.
- It has a single synchronization method which we will call *await()*.
- All P threads must call the *await()* method on the barrier.
 - In general the *await()* call *blocks*, and will not complete in *any* thread until the *last* of the P threads have made the call.
 - The *await()* call then completes in all P threads, and all can continue.



Life Thread with Barriers

```
Class LifeThread {  
    void run() {  
  
        while(true) {  
            for(int i = begin ; i < end ; i++)  
                for(int j = 0 ; j < N ; j++)  
                    sums [i] [j] = sum of cells values neighbouring (i, j) ;  
            barrier.await() ;  
            for(int i = begin ; i < end ; i++)  
                for(int j = 0 ; j < N ; j++)  
                    cells [i] [j] = update(cells [i] [j], sums [i] [j]) ;  
            barrier.await() ;  
        }  
    }  
}
```



Comments

- In this version, the first barrier synchronization call makes sure that no thread goes on to modify the `cells` array until all threads have finished using it's old value to calculate their elements of the `sums` array.
- The second barrier synchronization makes sure that no thread goes on to calculate the `sums` array in the `next` generation until all threads have finished calculating the new values of the `cells` array.



Being Safe

- If in doubt, put a barrier synchronization after each “decomposed loop” in the thread.
- Notice that with the barriers the behaviour of the program is now very similar to our first attempt, with separately created threads for each individual set of parallelized loops.
- **joining** with the main program followed by spawning new threads has a very similar effect to a barrier - but at greater cost!

Life in C, and OpenMP

- So long as we are using an appropriate version of C, pseudocode can look much the same:

```
while(1) {  
    for(int i = 0 ; i < N ; i++)  
        for(int j = 0 ; j < N ; j++)  
            sums [i] [j] = sum of cells values neighbouring (i, j) ;  
    for(int i = 0 ; i < N ; i++)  
        for(int j = 0 ; j < N ; j++)  
            cells [i] [j] = update(cells [i] [j], sums [i] [j]) ;  
}
```

- ... elided code for calculating sums and doing update also can be identical to Java.



OpenMP Parallel Version

```
#include <omp.h>
```

```
while(1) {
```

```
    #pragma omp parallel for
```

```
    for(int i = 0 ; i < N ; i++)
```

```
        for(int j = 0 ; j < N ; j++)
```

```
            sums [i] [j] = sum of cells values neighbouring (i, j) ;
```

```
    #pragma omp parallel for
```

```
    for(int i = 0 ; i < N ; i++)
```

```
        for(int j = 0 ; j < N ; j++)
```

```
            cells [i] [j] = update(cells [i] [j], sums [i] [j]) ;
```

```
}
```




Comments

- Impressively, this does *everything* our parallel Java version does - but *automatically*.
- It creates P threads, where by default P is the available number of cores.
- The `for(int i = 0 ; i < N ; i++)` loops after the *parallel for* pragma are automatically decomposed with an index subrange for each thread.
- By default a *barrier synchronization* is added after each parallel loop.



Summary

- Next week: Topics in *distributed memory and MPI intro*.



Further Reading

- See links embedded in these slides.