



4. Distributed Memory, and MPI Basics

Parallel Programming

Dr Hamidreza Khaleghzadeh
School of Computing
University of Portsmouth



Goals

- Set the background to *distributed memory programming* on clusters of computers.
- Basics of *MPI* - the parallel job.
- First example of message passing.



Limits of Shared Memory

- Previous lecture concentrated on parallel programming with *shared memory*.
- Has become important in recent years with rise of commodity *multicore* processors.
- But widely accepted there are limits to the *scalability* of shared memory systems - how many cores can share a common memory.
- In the search for scalability, we are forced to consider *distributed memory systems*.




Distributed Memory

- *Distributed Memory* (DM) systems are not particularly mysterious or unfamiliar.
 - A group of ordinary workstations connected by a LAN is perhaps the most basic instantiation.
- In supercomputers this idea is pushed to extremes, with thousands of *nodes* (essentially the workstations) connected by a special purpose high speed network or *interconnect*
 - e.g. *Myrinet*, *Infiniband*, or custom interconnects.

A Simple Cluster

- Consider a group of four similar workstations:



- Total memory available is four times that of individual workstation, but *distributed* over group.

A Continuum

- Can identify a whole range of systems that *could* be used for *distributed memory* parallel computation:

Ad hoc
grouping of
workstations
on TCP/IP
LAN in lab

*Clusters of
commodity
computers
dedicated to
parallel
computing
(e.g. Beowulf[†])*

Supercomputers
with thousands
of nodes on
custom
interconnect.



Increasing scale

[†] See for example http://en.wikipedia.org/wiki/Beowulf_cluster



From Threads to Processes

- From a software perspective, biggest change in programming these systems is that we move from *cooperating threads* - sharing access to a common memory - to *cooperating processes*, which each have their own *private memory*, but with no memory shared *between* processes.
- In our case these constraints are enforced by the physical hardware, but software considerations quite analogous to distinction between threads and processes within an individual *operating system*.



Programming

- *Programming* cooperating processes is generally harder than cooperating threads - often involves explicit *Inter-Process Communication* (IPC).
- For parallel processing at largest scales this seems hard to avoid.
- But quite similar programming techniques can be used over whole continuum of distributed memory systems introduced just now.



Message Passing

- In some sense the lowest common denominator approach, explicit Message Passing is *widely* used in DM parallel programs today.
- Many - probably most - parallel programs designed to run on large clusters and distributed memory supercomputers use one particular *API* (Application Programming Interface) for this.



The Message Passing Interface

- Some twenty-five years ago, *MPI version 1* was standardised by a group of manufacturers and academics called the *MPI Forum*.
 - Do *not* confuse *MPI* with *OpenMP*, which is a standard for *shared-memory* parallel programming we discussed last week[†]!
- API was defined for *C/C++* or *Fortran*, and still no *official* bindings beyond those two languages
 - though we will use an "unofficial" *Java* binding.

[†]To make things even more confusing, one of the popular open-source implementations of MPI is called *Open MPI*...



MPI BASICS



The MPI Programming Model

- MPI 1 assumes a fixed number (P , say) of processes participate in any given parallel program.
- This set of processes is called "*the world*"
 - Size of the world often but not always equal to number of physical nodes available[†].
- MPI implements the *SPMD model* - each process runs *the same* program, but of course operates on its own local memory (data).

[†]MPI 2 introduced the option to create new processes during program execution - "growing" the world. Unclear how widely this feature is used.



SPMD Programming

- *Single Program Multiple Data* in which all participant processors (or processes) run the same program image, but operate on their local memory contents.
- A special case of the more general *MIMD* (*Multiple Instruction Multiple Data*) model, in which different participants may run different local programs.
 - e.g. a common paradigm in *pre-MPI* days was for one node to run a *host* program that coordinated I/O and controlled the other nodes, which ran a separate *worker* programs that did most of the computation.



Process Rank

- MPI processes in the world of a particular parallel program execution are distinguished by rank
- Just the MPI name for the numeric identifier of the processor.
- If the size of the world is P , the rank runs from 0 to $P - 1$ [†]
- Directly comparable with `me` variable in our earlier parallel programming examples.

[†]Strictly this is rank relative to the world - ranks within *subgroups* of processes are also defined.



On Programming Examples

- We will generally use an *unofficial* Java API for MPI, as implemented in the software called *MPJ Express*.
- We will frequently refer to this Java API, *loosely*, as *MPJ*.
 - Strictly speaking, MPJ was another slightly different API[†], but since we will be using *MPJ Express* in labs, we better stick to the that variant of the API.
- In any case, this API adheres as closely as practical to the official *C* standard for MPI.

[†][MPJ: MPI-like Message Passing for Java](#)



MPJ Hello World

```
import mpi.* ;

public class HelloWorld {

    public static void main(String args[]) throws Exception {

        MPI.Init(args) ;

        int P    = MPI.COMM_WORLD.Size() ;
        int me = MPI.COMM_WORLD.Rank() ;

        System.out.println("Hi from <" + me + ">") ;

        MPI.Finalize() ;
    }
}
```




MPJ Fundamentals

- Every MPJ program is a *Java application* starting with a `main()` method.
- This application is started in *P* logical processes, typically running on different nodes of the system.
- Every MPJ program *must* begin by calling:

`MPI.Init(args) ;`

where `args` is the parameter of `main()`, and finish by calling:

`MPI.Finalize() ;`



The World

- The object:

`MPI.COMM_WORLD`

is an *MPI Communicator* spanning the world - the universe of processes in which the program is running.

- It is an instance of the class `Comm`, and in our binding many of the most useful methods in MPI belong to this class.
- For now, just accept that this is a Java object - and that most of the methods we see initially are called on this object.



Number of Processes and Process Rank

- In MPI, the number of processes the program runs in is *not* defined by the program itself.
- Instead it is defined by the command used to run the parallel program.
- The program can interrogate the number of processes by calling:
`MPI.COMM_WORLD.Size()`
- Each running process can find its “position” in the set of processes by calling:
`MPI.COMM_WORLD.Rank()`



Running MPJ HelloWorld

- At the command prompt on client computer, might run:

```
> mpjrun -np 4 HelloWorld
```

```
Hi from <1>
```

```
Hi from <3>
```


```
Hi from <2>
```

```
Hi from <0>
```

- Of course there is no particular order to the output from the 4 processes.


Scenario for Running an MPI Program

- Pool of available host computers, each running an *MPI service* or *MPI daemon*.



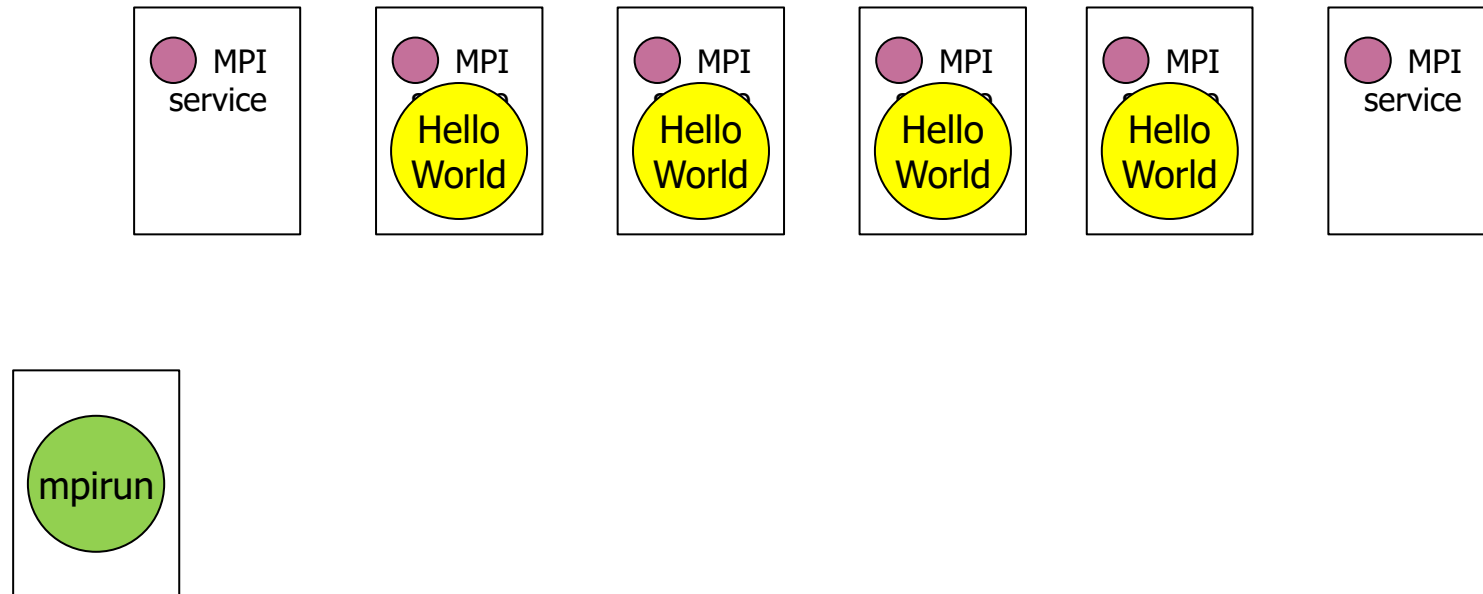
Scenario for Running an MPI Program

- Client program, e.g. `mpirun`, connects to P daemons and asks them to start processes.



Scenario for Running an MPI Program

- *Hello World* process (say) starts on P hosts.





Handling Print Statements

- Commonly, MPI daemons redirect any *standard output* of the *P* distributed processes (from *print* statements, etc) back to the client program - *mpirun* in our example
 - It is displayed on the console of the client process.
 - Useful for debugging, etc.
- Other I/O such as file input or output is *not* handled automatically - program has to figure out exactly where it wants it to go!



MESSAGE PASSING



Simple Message Passing


- The main point of MPI is to allow our P processes to *communicate* in order to collaborate on tasks.
- The simplest method for *sending* a message in MPJ is:
Send(buffer, offset, count, type, dest, tag)
- Unfortunately, the argument lists of operations in MPI tend to be long...



Message Buffer

- The *message buffer* is just an array in the user's program that contains the data to be sent to another program.
- The *buffer* argument of **Send()** is therefore a Java array containing this information.
- The *offset* and *count* arguments select the elements of *buffer* that are actually going to be sent in the message.
- The *type* argument defines the type of the elements of *buffer*. It is slightly redundant in Java - just follow the examples for now.

Buffer Example



- Elements of **buffer** actually sent are in **green**.
- Note **offset** is often 0.
- **count** may take value 1 to send a single element.



Destination and Tag

- The **dest** argument of **Send()** is just the numerical id - i.e. the *rank* - of the process to which the message is to be sent.
- The **tag** argument is a user-defined code for identifying the purpose of the message - for the time being we will just set it to 0.



Receiving a Message

- The simplest method for *receiving* a message in MPJ is:

Recv(buffer, offset, count, type, src, tag)

- Arguments are almost the same as **Send()**.
- **buffer** is now an array into which the contents of the incoming message will be placed; **offset** says where in the array, and **count** is the maximum message size that can be accepted.
- **src** is now the rank of the process we expect the message to come *from*.



An Example

- The following slide gives a sequential algorithm for calculating π , based on the following formula from Calculus:

$$\pi = \int_0^1 \frac{4}{1+x^2} dx$$

applying the "*rectangle rule*" for numerical integration (don't worry about the math!)

- N is the number of steps used for the numerical approximation.
- For clarity, we only give the central part of the algorithm.



Sequential π Calculation

```
double step = 1.0 / (double) N;  
double sum = 0.0;  
for(int i = 0 ; i < N ; i++) {  
    double x = (i + 0.5) * step ;  
    sum += 4.0 / (1.0 + x * x);  
}  
double pi = step * sum ;
```




MPJ Parallel π Calculation (Part 1)

```
int me = MPI.COMM_WORLD.Rank() ;
int P   = MPI.COMM_WORLD.Size() ;

int b = N / P ;
int begin = me * b ;
int end = begin + b ;


double step = 1.0 / (double) N ;
double sum = 0.0 ;
for(int i = begin ; i < end ; i++) {
    double x = (i + 0.5) * step ;
    sum += 4.0 / (1.0 + x * x);
}
```



Comments

- This is the easy part, and it is very much analogous to block-wise decompositions we have seen before in shared memory examples
- The number of processes, P , and the process identity me are now obtained by the MPJ inquiry functions $Size()$ and $Rank()$.

Communication Pattern Needed





MPJ Parallel π Calculation (Part 2)

```
if (me > 0) {  
    double [] sendBuf = new double [] {sum} ;    // 1-element array containing sum  
    MPI.COMM_WORLD.Send(sendBuf, 0, 1, MPI.DOUBLE, 0, 0) ;  
}  
else {    // me == 0 !  
    double [] recvBuf = new double [1] ;  
    for (int src = 1 ; src < P ; src++) {  
        MPI.COMM_WORLD.Recv(recvBuf, 0, 1, MPI.DOUBLE, src, 0) ;  
        sum += recvBuf [0] ;  
    }  
}  
double pi = step * sum ;
```



Comments

- This is rather harder to understand - there is a lot to digest here.
- Once absorbed, you should have a clearer idea of what SPMD programming with MPI is about!
 - Processes with $me > 0$ are sending a message containing their contribution to the sum to process 0.
 - Process 0 receives a message from processes 1 to $P-1$, and accumulates the value it receives into its contribution to the sum.
- Process 0 now holds the final result.



Summary

- Next week: *Communication in MPI*.