



## 6. Collective Communication

---

Parallel Programming

Dr Hamidreza Khaleghzadeh  
School of Computing  
University of Portsmouth



# Goals

---

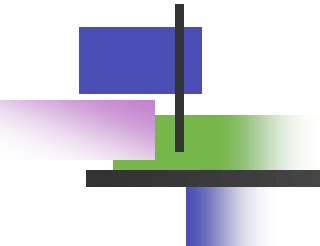
- Continue our survey of MPI by considering collective communications - an important part of the API.
- In several cases, motivate introduction of these methods by illustrating how their use can simplify parts of programs introduced in earlier labs.



# Collective Communication

---

- Past two lectures introduced fundamental methods for communication between processes of a distributed memory parallel program.
- These generally took the form of *point-to-point communication* - messages sent between *two* processes.
- A different and important paradigm (supported by MPI) involves *all* processes working together to move data between the memories of the processors.
- This is called *collective communication*.



---

# BROADCAST - THE SIMPLEST COLLECTIVE



# Idea of Broadcast

---

- The idea is fairly intuitive.
- One process needs to send a particular data item (embodied in MPI in an array) to *every other process* in the program.
- Very commonly the broadcasting process may be *process 0*, and the broadcast data may be some input data, or values that control the program as a whole
  - e.g. the size, *N*, of the problem to be solved, where this isn't declared globally as a constant.

# Broadcast using Send/Recv

```
int [] values = new int [M] ;
```

```
if (me == 0) {
```

```
    ... Initialize `values' - e.g. input them from user ...
```

```
    for (int dst = 1 ; dst < P ; dst++) {
```

```
        MPI.COMM_WORLD.Send(values, 0, M, MPI.INT, dst, 0) ;
```

```
    }
```

```
}
```

```
else {    // me > 0
```

```
    MPI.COMM_WORLD.Recv(values, 0, M, MPI.INT, 0, 0) ;
```

```
}
```

```
... Consistent values now available to all processes ...
```



# Critique

---

- This works fine, but requires the programmer to think about the details of how the broadcast is broken down into sends and receives.
- Moreover this implementation is *less efficient* than it needs to be - it probably takes time  $O(M \times P)$  to complete.
  - There are much more efficient algorithms to implement a broadcast, but their logic is more complex.



# Broadcast using a Collective

---

```
int [] values = new int [M] ;
```

```
if (me == 0) {
```

```
    ... Initialize `values' - e.g. input them from user ...
```

```
}
```

```
MPI.COMM_WORLD.Bcast(values, 0, M, MPI.INT, 0) ;
```

```
... Consistent values now available to all processes ...
```





# The Bcast Method

---

- The new method looks like this:

**Bcast**(buffer, offset, count, type, root)

where **buffer**, **offset**, **count** and **type** describe source and destination arrays (as previously) and **root** is rank of broadcasting process.

- Importantly, **Bcast** must be called by *all processes*, “at the same point” in a program, and with *consistent arguments* (e.g. all must agree on the values of **root**, **type**, etc).
  - Recalls usage of *barrier* in shared memory programs.



# Advantages of Bcast

---

- The user code is shorter, because the logic of the communication pattern is captured in the library.
- Perhaps more importantly it should be faster - a well-tuned implementation of Bcast may complete in time  $O(M + \log(P))$  for large messages and numbers of processors.
- Broadcast is an archetype for a whole family of collective communications.



# REDUCTION



# Reduce vs Broadcast

---

- *Reduction* is in a sense the opposite operation to broadcast.
- Here values are taken from *all* processes, reduced to *single values* by some *combining operation*, and those single values are deposited on a "root" process.
- We have already seen this kind of pattern in our first MPJ code for calculating  $\pi$ , where the combining operation was floating point addition.



# MPJ Parallel $\pi$ Collecting Results

---

```
if (me > 0) {
    double [] sendBuf = new double [] {sum} ;
    // 1-element array containing sum

    MPI.COMM_WORLD.Send(sendBuf, 0, 1, MPI.DOUBLE, 0, 0) ;
}

else {    // me == 0 !

    double [] recvBuf = new double [1] ;

    for (int src = 1 ; src < P ; src++) {
        MPI.COMM_WORLD.Recv(recvBuf, 0, 1, MPI.DOUBLE, src, 0) ;
        sum += recvBuf [0] ;
    }
}

double pi = step * sum ;
```



# Collecting $\pi$ Results using Reduce

---

```
double [] sendBuf = new double [] {sum} ;  
    // 1-element array containing local sum  
  
double [] recvBuf = new double [1] ;  
  
MPI.COMM_WORLD.Reduce(sendBuf, 0, recvBuf, 0,  
    1, MPI.DOUBLE, MPI.SUM, 0) ;  
  
if (me == 0) {  
    double pi = step * recvBuf [0] ;  
    ...  
}
```



# Reduce Interface

---

- MPJ interface looks like:

**Reduce**(sendbuf, sendoffset, recvbuf, recvoffset,  
count, type, op, root)

- If **count** is  $> 1$ , **sendbuf** arrays from  $P$  processes are combined element by element to produce an array of **count** results in **recvbuf**.
- **op** is the combining operation, and it can take following values:
  - **MPI.SUM, MPI.PROD, MPI.MAX, MPI.MIN, MPI.LAND, MPI.BAND, MPI.LOR, MPI.BOR, MPI.LXOR, MPI.BXOR, MPI.MINLOC** and **MPI.MAXLOC**



---

# GATHER AND SCATTER



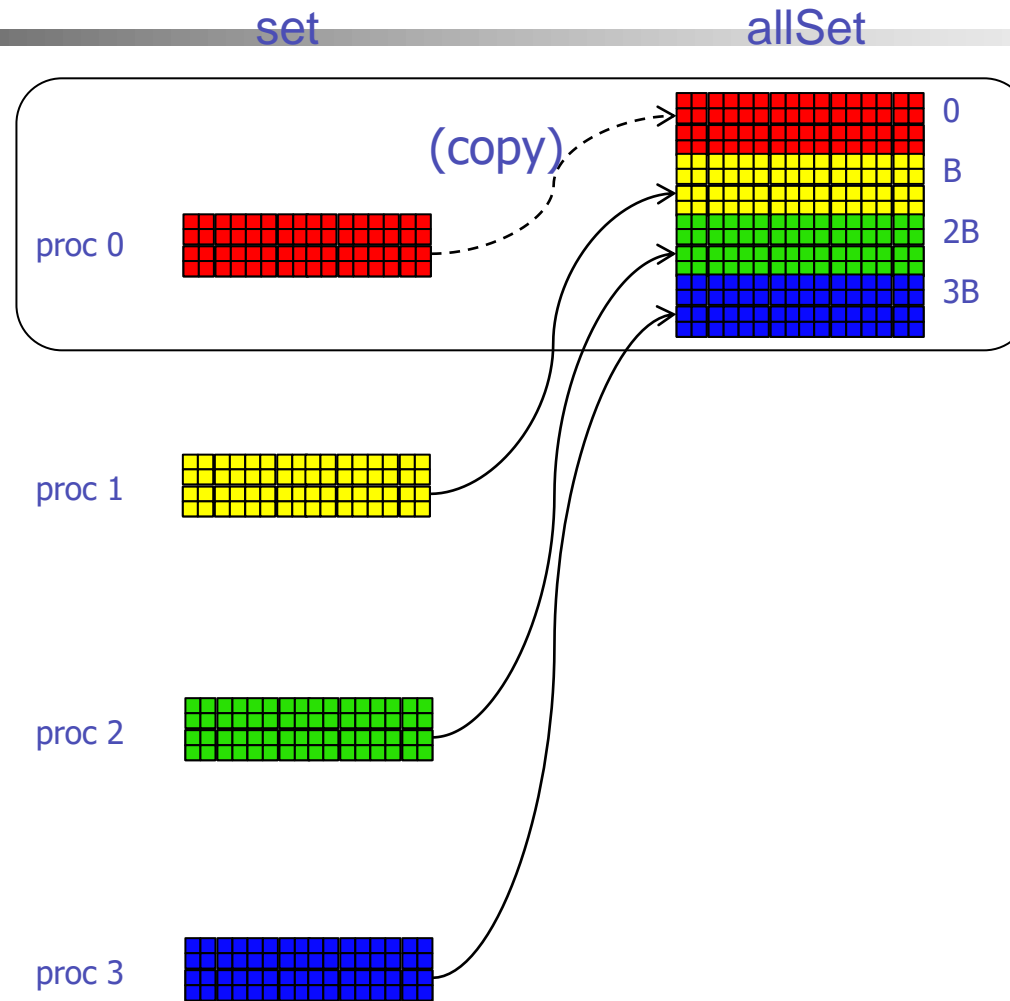


# MPJ Mandelbrot Revisited

---

- In the week 5 lecture, we sketched one possible implementation of a Mandelbrot Set calculation using MPJ.
- The only non-trivial communication required there was collecting together results at the end of the calculation, reproduced below.

# Communication Pattern



# Gathering Results in Mandelbrot

```
if(me > 0) {
    MPI.COMM_WORLD.Send(set, 0, B, MPI.OBJECT, 0, 0);
}
else { // me == 0
    for(int i = 0 ; i < B ; i++) { // copy local `set' to start of `allSet'
        for(int j = 0 ; j < N ; j++) {
            allSet [i] [j] = set [i] [j] ;
        }
    }
    for(int src = 1 ; src < P ; src++) {
        MPI.COMM_WORLD.Recv(allSet, src * B, B, MPI.OBJECT, src, 0);
    }
    ... display allSet ...
}
```



# The "gather" operation

---

- This behaviour is captured in the general "gather" operation - one of the collective operations supported directly in MPI.
- In MPJ the interface is fairly complex because it involves two separate arrays:

**Gather**(sendbuf, sendoffset, sendcount, sendtype,  
recvbuf, recvoffset, recvcount, recvtype,  
root)



# Results in Mandelbrot using Gather

---

```
MPI.COMM_WORLD.Gather(set, 0, B, MPI.OBJECT,  
                      allSet, 0, B, MPI.OBJECT,  
                      0) ;
```

```
if(me == 0) {  
    ... display allSet ...  
}
```

- Results sent from  $P$  processes automatically get concatenated together into  $P \times B$  elements of the `recvbuf` array, where  $B$  is the value of `recvcount` argument.
- In obscure cases `sendcount` and `recvcount` could be different, or `sendtype` and `recvtype` could be different. Usually, as here, they take the same values.

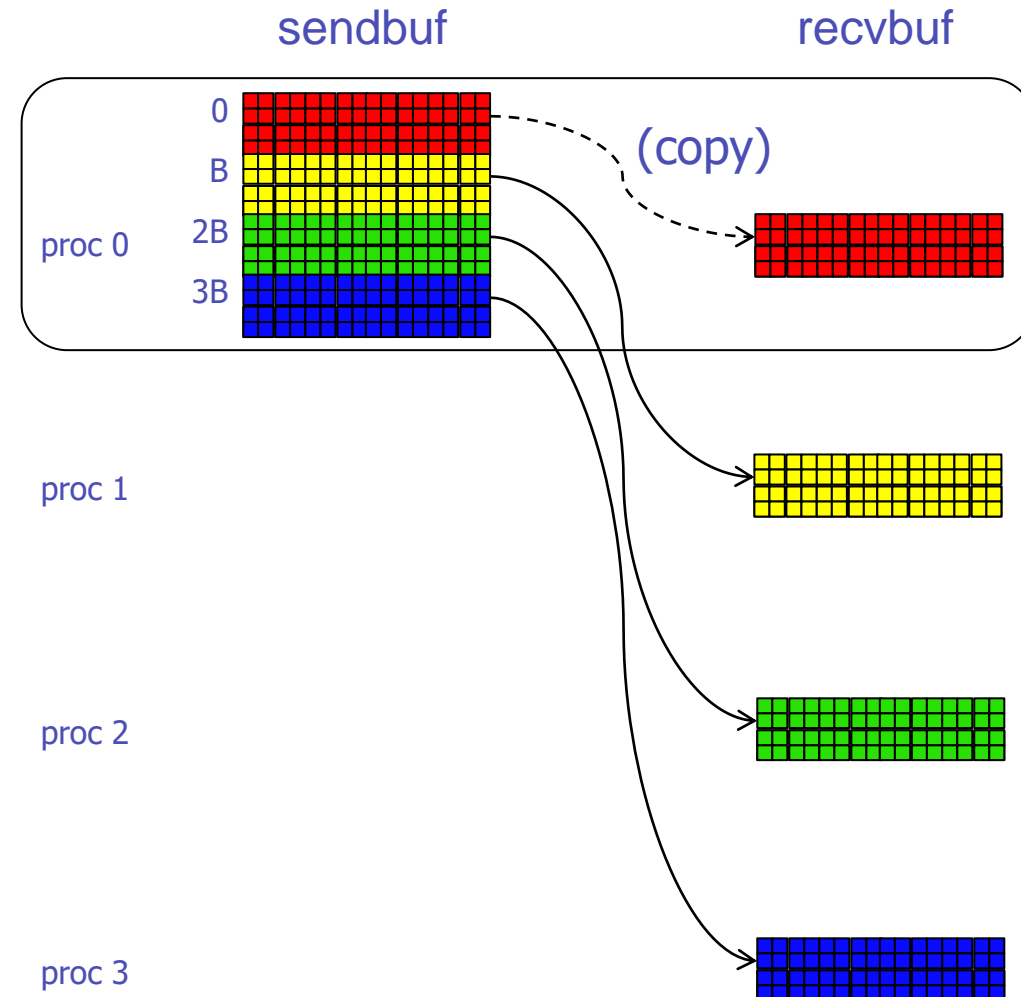


# Scatter

---

- *Scatter* is the opposite operation to gather.
- A common scenario is where process 0 (say) initializes the values in a large array (e.g. by reading from a file), then these values have to be distributed across all processes for data-parallel processing.
- In MPJ, argument list of *Scatter* is identical to *Gather*, but now the first argument *sendbuf* is the “large” array that will be divided up into *B*-sized chunks.

# Scatter Communication Pattern



(Assume root is 0.)

(Again assume 2d arrays  
– not necessary!)



# Allgather


---

- The MPI *Allgather* operation behaves like a gather followed by a broadcast.
- As a motivating example, consider the simulation of  $N$  stars in a galaxy moving under the force of gravity.
- Each individual star feels the gravitational force of every other star, according to *Newton's inverse square* law of gravity.



# Forces on stars

- Each star (one selected here) feels force of gravity from every other star.





# Total Force on a star

---

- Force on the  $i$ 'th star:

$$\mathbf{F}_i = \sum_{j \neq i} \frac{G \mathbf{n}_{ij}}{(\mathbf{r}_i - \mathbf{r}_j)^2}$$

- Here  $\mathbf{r}_i$  is (3-vector) position of  $i$ 'th star and  $\mathbf{n}_{ij}$  is a vector that takes into account the direction of the individual force (along the line from  $i$  to  $j$ ).
- Don't worry about mathematical details! I have assumed all stars have same mass.

# Decomposition of Star Positions

proc 0



B

proc 1



B

proc 2



B

proc 3



B

- Can assume this is array of x-positions.
- Actually need two more arrays for y and z coordinates (and perhaps three more for components of velocity).




# Calculating Forces

---

- Each process updates positions and velocities of  $B$  stars.
- But when we come to calculate accelerations, need to know positions of every other star.
- Can use **AllGather** - which effectively combines a call to **Gather** with a call to **Bcast** - to get current position of every star to every process.
- Each process then does  $N \times B$  individual force calculations.

# AllGather Communication Pattern





# AllReduce

---

- Finally, another equally useful “combined” operation is *AllReduce*, which behaves like a *Reduce* followed by a broadcast.
- For example, an alternative implementation of the stars simulation may see all star positions held *replicated* across all nodes.
- A node calculates the net force a block of these stars exert on on *all* other stars, then perform an *AllReduce* on the resulting array.
- All nodes redundantly perform the less demanding position/velocity updates.



# Summary

---

- Finished our discussion of the MPI standard with an exploration of *collective communications*, which in MPJ are defined in the *Intracomm* class.
- Standardizing these was a significant contribution of MPI, and they are a widely used part of the API.
- Discussion hasn't been exhaustive, and in fact the general idea of collective communication can be extended well beyond what is captured in MPI.



# Further Reading

---

- MPJ Express API:

<http://mpj-express.org/docs/javadocs/index.html>

- William Gropp, Ewing Lusk and Anthony Skjellum, *Using MPI*, 2<sup>nd</sup> Edition MIT Press, 1999.
  - Standard text on MPI, but examples are in C and Fortran.
  - Available as an electronic book through the library.