



## 10. Cilk and Related Systems

---

Parallel Programming

Dr Hamidreza Khaleghzadeh  
School of Computing  
University of Portsmouth



# Goals

---

- Introduce the *Cilk* parallel programming system and its programming model.
- Illustrate with simple examples like *QuickSort*
- Briefly review the related *Fork/Join framework* in Java 1.7, and *Intel Threading Building Blocks*.
- A glance at another recursive algorithm - Barnes Hut.



# Cilk

---

- *Cilk* is a parallel programming language developed at MIT in the mid-1990s by Charles Leiserson and collaborators.
- It is particularly well suited to a *divide-and-conquer* (recursive) style of parallelism.
- Works by repeatedly spawning *lightweight threads* that are *micro-scheduled* by the run-time system of the language.
- The scheduling mechanism built into the language is called *work-stealing*.



# Cilk Developments

---

- The original *Cilk* is a small extension to *ANSI C*.
- In 2006 a version of Cilk was commercialized by a spin-off from MIT called Cilk Arts.
- Acquired by Intel around 2009, and now marketed as *Cilk Plus* - an extension of *C++*.
- Ideas similar to Cilk are now incorporated in other systems - e.g. the *Fork/Join framework* in Java 1.7.



# The Fibonacci Series

---

- Fibonacci sequence probably familiar:  
0, 1, 1, 2, 3, 5, 8, 13, 21, ...,  $\text{Fib}_n$ , ...
- Each term is the sum of the previous two terms.
- There is a recursive calculation of the  $n$ th term in this series, beloved of functional programmers (despite or because of it being a spectacularly inefficient algorithm).
- A standard C version of this algorithm is given on the next slide.



# Fibonacci Function (C)

---

```
int fib (int n) {  
    if (n<2) {  
        return n;  
    }  
    else {  
        int x, y;  
        x = fib (n-1);  
        y = fib (n-2);  
        return (x+y);  
    }  
}
```



# Divide and Conquer

---

- This recursive pattern, where a calculation is broken into two or more smaller calculations, is often called *divide and conquer*.
- It can often give rise to a natural form of parallelism in which the *sub-tasks are run concurrently*.
- Cilk can be particularly effective for expressing this kind of parallelism.



# Fibonacci Function (Cilk)

---

```
cilk int fib (int n) {  
    if (n<2) {  
        return n;  
    }  
    else {  
        int x, y;  
        x = spawn fib (n-1);  
        y = spawn fib (n-2);  
        sync ;  
        return (x+y);  
    }  
}
```






# New Keywords

---

- The core of Cilk uses just three new keywords: *cilk*, *spawn* and *sync*.
  - Keyword *cilk* identifies a *cilk procedure* - a parallel version of a C function.
  - Keyword *spawn* indicates a cilk procedure call be run *in parallel* with the calling thread.
  - Return values *written* by spawned procedures cannot be safely read until a *sync* statement is executed by the calling procedure - a kind of barrier or join synchronization across threads spawned by this procedure.

# Thread Creation in fib





# Proliferation of Threads

---

- Very rapidly have *many* (in this case exponentially many) lightweight threads.
- But Cilk micro-scheduling is *very efficient*, and overheads of *spawn/sync* operations is small.
  - At least an order of magnitude less than if these were, say POSIX threads or Java threads<sup>†</sup>
- Number of underlying POSIX-style threads typically equal to number of cores
  - execution of the lightweight Cilk threads shared between these.

<sup>†</sup>For the similar TBB (see later) thread vs task overhead ratio is 18 for Linux, 100 for Windows, according to TBB User Guide from Intel.

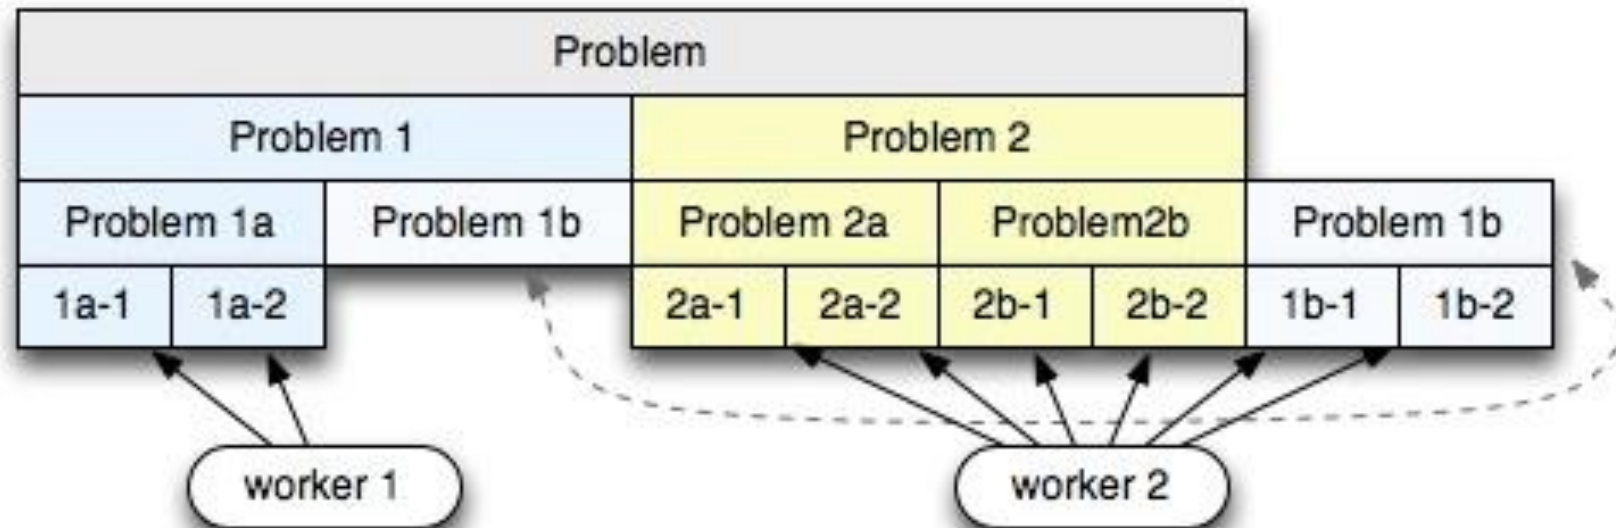


# Work Stealing

---

- Each core maintains a *queue* (specifically a double-ended queue, or *deque*) of pending cilk procedures ("threads") awaiting execution.
- When one core runs out of work, it randomly chooses one of the other cores, and *steals* a pending procedure from the top of the "victim's" work deque.
  - "Near the top" implies this is a task spawned early in the recursion, which will probably spawn further threads as the "thief" processes it.
  - Minimizes synchronization between cores.

# Visualization of Work Stealing<sup>†</sup>



<sup>†</sup>Image from

<http://www.igvita.com/2012/02/29/work-stealing-and-recursive-partitioning-with-fork-join/>



# A More Interesting Example

---

- In the *QuickSort* algorithm, an array or section of an array is sorted by first *partitioning* it so that elements on the *left* side of the array are *less than* some randomly chosen *pivot value*, and those on the *right* side are *greater than* the pivot.
- The two sides of the partition are then independently sorted by a recursive call.
- A C version of the core QuickSort algorithm is given on the next slide.



# QuickSort in C

---


```
void quicksort(int data [], int first, int last) {  
    if(last > first) {  
        int pivotIndex = last ;  
        pivotIndex = partition(data, first, last, pivotIndex) ;  
        quicksort(data, first, pivotIndex - 1) ;  
        quicksort(data, pivotIndex + 1, last) ;  
    }  
}
```

# Behaviour of partition()

- Before call to partition:



- After call to partition:







# Notes

---

- Both **quicksort** and **partition** functions operate on a subset of the **data** array determined by the index values **first** and **last**.
- The **partition** function can be implemented in various ways using  $O(N)$  comparisons and swaps of pairs of elements.
  - One particular algorithm is given at <http://en.wikipedia.org/wiki/Quicksort>
- Note the ***divide and conquer*** nature of this very widely used algorithm.



# QuickSort in Cilk

---

```
cilk void quicksort(int data [], int first, int last) {  
    if(last > first) {  
        int pivotIndex = last ;  
        pivotIndex = partition(data, first, last, pivotIndex) ;  
        spawn quicksort(data, first, pivotIndex - 1) ;  
        spawn quicksort(data, pivotIndex + 1, last) ;  
        sync ;  
    }  
}
```



# Notes

---

- Simply change `quicksort` to a `cilk` procedure, and `spawn` it rather than calling it (remember to `sync`!)
- The most intricate part of the algorithm - the in-place `partition` function - hasn't changed at all. It remains an ordinary C function.
- For optimal results, parallel divide and conquer usually sets some *threshold problem size* below which it reverts to sequential computation.
  - I haven't done this here - we `spawn` all the way down to the *base case* of the recursion.



# QuickSort Performance

---

- **C version:**

```
$ gcc -O2 qsort.c -o qsort
```

```
$ ./qsort
```

Completed in 1376 milliseconds

- **Cilk version:**

```
$ cilkc -O2 qsort.cilk -o qsort
```

```
$ ./qsort
```

Completed in 1547 milliseconds

```
$ ./qsort --nproc 2
```

Completed in 865 milliseconds



# Notes on Timings

---

- Timings on my (dual core) laptop, for  $N = 10,000,000$
- Overhead of converting C to Cilk here only around 12% - could probably be reduced by "thresholding" as described earlier.
- Speedup of Cilk version on two cores close to 1.8.
- -O2 is a recommended optimization level (without it C and Cilk both 2-3 times slower)



# Postscript

---

- Sadly last MIT release of original Cilk was in 2007, and, while still available on their web site, no longer seems to be maintained.
- *Cilk Plus* keywords (`cilk_spawn`, `cilk_sync`) still available in Intel C++ compiler, but deprecated around 2017.
  - Intel now recommended to use *OpenMP* pragmas `task` and `taskwait` to replace these keywords. See:

<https://software.intel.com/en-us/articles/migrate-your-application-to-use-openmp-or-intelr-tbb-instead-of-intelr-cilktm-plus>

- Possible rebirth at MIP?: <http://cilk.mit.edu/>



# Work Stealing in Java 1.7

---

- Java 1.7 introduced an implementation of fine-grained “threads” similar to that in Cilk, with its Fork/Join Framework.
- Basic ideas proposed by Doug Lea around 2000, and explicitly inspired by Cilk.



# QuickSort using Fork/Join

```
public class ForkJoinQuickSort extends RecursiveAction {  
    ...  
    int first, last ;  
    ForkJoinQuickSort(int first, int last) {  
        this.first = first ;  
        this.last = last ;  
    }  
    protected void compute() {  
        if(last > first) {  
            int pivotIndex = last ;  
            pivotIndex = partition(data, first, last, pivotIndex) ;  
            invokeAll(new ForkJoinQuickSort(first, pivotIndex - 1),  
                    new ForkJoinQuickSort(pivotIndex + 1, last)) ;  
        }  
    }  
}
```





# Notes

---

- The `main` method does this:

```
ForkJoinQuickSort mainTask =  
    new ForkJoinQuickSort(0, N-1) ;  
ForkJoinPool mainPool = new ForkJoinPool(P) ;  
mainPool.invoke(mainTask) ;
```

- `compute` method similar to `run` method on `Thread`.
- My implementation assumed `data` is a static field of the class `ForkJoinQuickSort`.



# Java Fork/Join Performance

---

- Timings on my laptop...
- With  $P = 1$ :
  - \$ java ForkJoinQuickSort
  - Completed in 2818 milliseconds
- With  $P = 2$ :
  - \$ java ForkJoinQuickSort
  - Completed in 1544 milliseconds
- Speed up on 2 cores about 1.8 again.
- Slower than Cilk, but not optimized - thresholds may be more important here?



# Intel Threading Building Blocks

---

- Intel *Threading Building Blocks* contain some ideas on task parallelism similar to Cilk, etc.
- A C++ template library, rather than C language extensions or Java classes.
- Task spawning features quite similar to Java 7 F/J, but lower-level
  - Note TBB predates release of Java F/J Framework by several years.
- Users encouraged, rather, to use higher level, more “declarative” templates from TBB, like `parallel_for`, `parallel_reduce`, ...



```
class FibTask: public tbb::task {
```

```
public:
```

```
    const long n;
```

```
    long* const sum;
```

```
    FibTask( long n_, long* sum_ ) : n(n_), sum(sum_) {}
```

```
    task* execute() {
```

```
        if(n < 2) {
```

```
            *sum = n ;
```

```
        }
```

```
        else {
```

```
            long x, y;
```

```
            FibTask& a = *new( allocate_child() ) FibTask(n-1,&x);
```

```
            FibTask& b = *new( allocate_child() ) FibTask(n-2,&y);
```

```
            // Set ref_count to 'two children plus one for the wait'.
```

```
            set_ref_count(3);
```

```
            spawn( b );
```

```
            spawn_and_wait_for_all(a);
```

```
            *sum = x+y;
```

```
        }
```

```
        return NULL;
```

```
    }
```

```
};
```



# Task Parallel Fibonacci in TBB

---

- Sample on previous slide adapted from TBB User Guide.
- Main program might do:

```
long sum;
```

```
FibTask& a = *new(tbb::task::allocate_root()) FibTask(N, &sum);
```

```
tbb::task::spawn_root_and_wait(a) ;
```

```
printf("sum = %ld\n", sum) ;
```

- Not all TBB is so gruesome, but it does require some expertise in C++!



---

Special Topic:

# **BARNES-HUT N-BODY CODE**



# N-Body Codes

---

- *N-Body codes* are codes that simulate the movement of “*N*” bodies (*N* usually large) under some law of mutual attraction or repulsion.
- Examples include the “Galaxy” simulation and the molecular dynamics simulation posted as development project ideas in the week 8 lab script.
- Where the force law has a long range nature - e.g. the inverse square law of gravity - the Barnes-Hut method may be effective.



# Complexity of Naïve Algorithm

---

- The “naïve” algorithm for calculating force on all particles involves  $N$  force law contributions on each of  $N$  particles - complexity  $N^2$ .
- In Barnes-Hut we group together particles exerting a force on  $i$ th particle.
- If group is far from particle  $i$ , approximate force exerted by group as a whole as force exerted by a *single large mass* at the *centre of gravity* of the group.







# Grouping: the Barnes Hut Tree

---

- First divide cubical region of 3d space into  $2^3 = 8$  regions, halving each dimension.
- For every sub-region that contains any particles, divide again recursively to make an *oct-tree*, until “leaf” regions have at most one particle.

# Two Dimensional Example


Adaptive quadtree where no square contains more than 1 particle



- Picture borrowed from [www.cs.berkeley.edu/~demmel/cs267/lecture26/lecture26.html](http://www.cs.berkeley.edu/~demmel/cs267/lecture26/lecture26.html)


# Barnes Hut Force Computation

- Force on a particle  $i$  - start with root node  $n$ :
  - if node  $n$  is "distant from"  $i$ , just add contribution to force on  $i$  from centre of mass of  $n$  - no need to visit children of  $n$ ;
  - if node  $n$  is "close to"  $i$ , visit children of  $n$  and recurse.
- Hinges on definition of *distant from* and *close to*.
  - Basic idea is that a node representing some region of space is *distant* from a particle  $i$  if the angle it subtends is smaller than a threshold *opening angle*:



# Complexity

- On average, number of nodes “opened” to compute force on  $i$  is  $O(\log N)$ , as opposed to visiting  $O(N)$  particles in naïve algorithm.
- Huge win when  $N \approx 10^{10}$ , for example (c.f. week 1).



Barnes-Hut version of my  
“Galaxy” program, with  
100,000 stars.



## Scope for Recursive Parallelization?

---

- *Idea*: can we exploit the kind of recursive parallelization approaches discussed in this lecture to parallelize Barnes-Hut?
- I will make a sequential Java version of B-H available on Moodle.
- The core `calcForce()` method on the `Node` class has a recursive structure similar to other examples discussed in this lecture - maybe can use Java Fork/Join or similar to parallelize?



# Summary

---

- Reviewed Cilk, one of the more elegant and efficient systems for shared memory parallel programming in C.
- Especially suited for *divide and conquer parallelism* (while OpenMP is arguably suitable for *data parallelism*).
- With emergence of multicore, Cilk has been commercialized, and its ideas continue to influence other products and systems.
- Also briefly reviewed products influenced by Cilk - Intel TTB, and Java Fork/Join.



# Further Reading

---

- *Cilk 5.4.6 Reference Manual*, Supercomputing Technologies Group, MIT Laboratory for Computer Science, <http://supertech.csail.mit.edu/cilk/>
- Doug Lea, *A Java Fork/Join Framework*, ACM 2000 Java Grande Conference, <http://gee.cs.oswego.edu/dl/papers/fj.pdf>
- JDK 1.7 Fork/Join Framework, <http://docs.oracle.com/javase/tutorial/essential/concurrency/forkjoin.html>